

NVIDIA STOCK PREDICTING-ARIMA

Members : Daniel, Brinta



OpenAI | NVIDIA

Content *overview*

1

DATA INTRO

2

ARIMA INTRO

3

CLEANING PROCESS

4

MODEL BUILT

5

RESULT

6

REFERENCES

Data Intro (Why Nvidia) ?

NVIDIA :

- **Founded:** 1993
- **Founder & CEO:** Jensen Huang
- **Headquarters:** Santa Clara, California
- **Core Businesses:** GPUs, AI computing, data centers, autonomous vehicles, cloud infrastructure

Market Significance

- **Ticker Symbol:** NVDA (NASDAQ)
- **Current Stock Price :** ~\$190/share - Start from \$5 (2017)
- Became one of the **top 5 most valuable U.S. companies**

Why NVIDIA Was Selected

- High volatility makes it ideal for short-term forecasting
- Stock behavior reflects real-world macro trends (tech cycles, AI investments, etc.)




NasdaqGS - BOATS Real Time Price • USD

NVIDIA Corporation (NVDA) ⚡ Follow ⚡ Analyze with AI

190.17 +3.31 +(1.77%) 192.25 +2.08 (+1.09%)

At close: November 14 at 4:00:01 PM EST Overnight: 8:58:34 PM EST ⓘ



NVIDIA Stock Dataset & Variable Breakdown

Dataset Source

- Provider: NVIDIA Historical Stock Price Dataset (Yfinance)
- Time Range: 2 years of data
- Total Observations: ~500 rows (daily frequency)
- Target Variable: **Close price**

Price	Close	High	Low	Open	Volume
Ticker	NVDA	NVDA	NVDA	NVDA	NVDA
count	1255.000000	1255.000000	1255.000000	1255.000000	1.255000e+03
mean	62.886648	63.968154	61.690980	62.891741	3.969987e+08
std	54.963195	55.806759	54.060165	55.027769	1.819776e+08
min	11.213526	11.720918	10.800024	10.957835	9.788400e+07
25%	18.584849	18.916072	18.067006	18.413221	2.460908e+08
50%	31.892086	32.694478	31.287299	32.217442	3.824624e+08
75%	114.261547	116.566815	111.538470	114.100646	5.072100e+08
max	207.039993	212.190002	205.559998	208.080002	1.543911e+09

Data Overview

The dataset provides the following columns for analysis:

- Date**: The date of the stock data entry.
- Price**: The price of the stock on that date.
- Adj Close**: The adjusted closing price, accounting for stock splits and dividends.
- Close**: The closing price of the stock on that date.
- High**: The highest price during the trading day.
- Low**: The lowest price during the trading day.
- Open**: The opening price on the day.
- Volume**: The number of shares traded on that day.

```
df = pd.DataFrame(data)
display(df.head(5))
```


	Price	Close	High	Low	Open	Volume
Ticker	NVDA	NVDA	NVDA	NVDA	NVDA	NVDA
Date						
2020-11-16	13.474119	13.607960	13.115962	13.132412	413776000	
2020-11-17	13.381402	13.554872	13.263263	13.511256	312028000	
2020-11-18	13.387879	13.564341	13.144872	13.424268	510924000	
2020-11-19	13.399345	13.446451	13.060131	13.172538	565936000	
2020-11-20	13.047923	13.453435	13.025241	13.413059	341088000	

What Is an ARIMA Time Series Model?

- **ARIMA** = AutoRegressive Integrated Moving Average
 - **AR:** Past values to predict future values.
 - **I:** Remove trends and make the series stationary.
 - **MA:** Uses past forecasting errors -> improve predictions.
- Why It's Useful:
 - Captures trends + patterns over time
 - Works well for finance, sales, demand forecasting
 - **short-term forecasting**

ARIMA Equations

$$\left(1 - \sum_{i=1}^{p'} \alpha_i L^i\right) X_t = \left(1 + \sum_{i=1}^q \theta_i L^i\right) \epsilon_t$$






Autoregressive
Integrated Moving
Average Prediction
Model

[ə-tō-ri-ri-'gre-siv 'in-tə-grā-təd
'mü-vij 'a-v(a-)rij pri-'dik-shən
'mä-dəl]

A statistical analysis model
that uses time series data
to predict future trends.

Stationarity Test- ADF test





DATA CLEANING

```
data.isna().sum()
```

0

Price Ticker

Close	NVDA	0
High	NVDA	0
Low	NVDA	0
Open	NVDA	0
Volume	NVDA	0

dtype: int64


```
# Extract the 'Close' price as the target variable  
target = data['Close']  
  
# Drop the 'Close' column from the original DataFrame  
data = data.drop('Close', axis=1)  
  
# Display the first few values of the target variable  
print("\nTarget variable (Close Price):")  
print(target.head())
```

Target variable (Close Price):
Ticker NVDA
Date
2023-11-27 48.213585
2023-11-28 47.792831
2023-11-29 48.111649
2023-11-30 46.742455
2023-12-01 46.737453

```
print("\nModified DataFrame (data) after dropping 'Close':")  
print(data.tail())
```

Modified DataFrame (data) after dropping 'Close':

Price	High	Low	Open	Volume
Ticker	NVDA	NVDA	NVDA	NVDA
Date				
2025-11-20	196.000000	179.850006	195.949997	343504800
2025-11-21	184.559998	172.929993	181.240005	346926200
2025-11-24	183.500000	176.479996	179.490005	256618300
2025-11-25	178.160004	169.550003	174.910004	320600300
2025-11-26	182.910004	178.240005	181.630005	183852000




Test if NVIDIA stock price series is stationary or not?

Stock Price Over time :

- Rolling mean and rolling standard deviation change over time
- Strong upward trend → non-stationary behavior
- The p-value is $0.99 \gg 0.05$
- fail to reject the null hypothesis
- Conclusion : the series is non-stationary

Action :

- need to transform to stationary form before modeling



Results of dickey fuller test

Test Statistics	-1.514419
p-value	0.526367
No. of lags used	4.000000
Number of observations used	498.000000
critical value (1%)	-3.443549
critical value (5%)	-2.867361
critical value (10%)	-2.569870

dtype: float64


```
from statsmodels.tsa.stattools import adfuller
```

```
result = adfuller(target)
print('ADF Statistic: ', result[0])
print('p-value:', result[1])
```

```
ADF Statistic: -1.5144190504473505
p-value: 0.5263667651818924
```

Why Log Transform the Data

- **Stabilize Volatility:** It ensures that price swings (volatility) are consistent over time, regardless of the stock's absolute price.
- **Normalize Scale:** It converts exponential price growth into a more linear trend, which is better suited for linear models like ARIMA.
- **Focus on Returns:** It allows the model to analyze percentage changes (returns), which are more relevant in finance than absolute dollar changes.
- **Meet ARIMA Assumptions:** It helps satisfy the statistical assumption that the data's variance remains constant Homoscedasticity.



Understanding the Price Distribution

1. Price Density Plot



- Shows what price has occurred more frequently
- Bimodal curve - not normally distributed
- 1st big peak around 120\$-150\$
- 2nd smaller peak at 170\$-200\$- showing the recent growth and AI boom

2. Trend & Seasonal Pattern

The graph is broken down into:

- Trend: long-term direction (upward)
- Seasonal: Spike and trough pattern
- Fluctuates between 1.00 to 1.02- typically add or subtract about 2% from the price
- Noise: random fluctuations around 1.00 mark

```
decomposition = seasonal_decompose(df['Close'], model='additive', period=30)
result = seasonal_decompose(df['Close'], model='multiplicative', period = 30)
fig = plt.figure()
fig = result.plot()
fig.set_size_inches(16, 9)
```



Preparing the Time Series for ARIMA


Why is it important :

- ARIMA only works on data that is stationary
- Stationary = no trend, no changing variance, no seasonality
- The raw NVIDIA price was not stationary
(p-value = 0.99)

```
#taking log transformation of closing price
from pylab import rcParams
rcParams['figure.figsize'] = 10, 6
df_log = np.log(df_close)
moving_avg = df_log.rolling(12).mean()
std_dev = df_log.rolling(12).std()
plt.legend(loc='best')
plt.title('Moving Average')
plt.plot(std_dev, color ="black", label = "Standard Deviation")
plt.plot(moving_avg, color="red", label = "Mean")
plt.legend()
plt.show()
```

What We Did

- Applied log transform to reduce volatility
- Used rolling mean visualization to inspect trend
- Stabilized variance and reduced trend effect



DEVELOP ARIMA MODEL


ARIMA :

- Split to test the model on unseen future data
- Prevents overfitting
- train on history → predict future

After log transformation, the data was divided into:

- 90% Training Data (fit the ARIMA model)
- 10% Testing Data (used to evaluate predictions)

```
#split data into train and testing set
train_data, test_data = df_log[3:int(len(df_log)*0.9)], df_log[int(len(df_log)*0.9):]
plt.figure(figsize=(10,6))
plt.grid(True)
plt.xlabel('Dates')
plt.ylabel('Closing Prices')
plt.plot(df_log, 'green', label='Train data')
plt.plot(test_data, 'blue', label='Test data')
plt.legend()
```




ARIMA Model Selection & Results

ARIMA:

```
Performing stepwise search to minimize aic
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-1785.697, Time=0.15 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=-1788.641, Time=0.09 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-1787.935, Time=0.18 sec
ARIMA(0,1,0)(0,0,0)[0] : AIC=-1784.107, Time=0.07 sec
ARIMA(2,1,0)(0,0,0)[0] intercept : AIC=-1789.952, Time=0.16 sec
ARIMA(3,1,0)(0,0,0)[0] intercept : AIC=-1790.270, Time=0.16 sec
ARIMA(3,1,1)(0,0,0)[0] intercept : AIC=-1789.451, Time=1.07 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=-1788.811, Time=0.40 sec
ARIMA(3,1,0)(0,0,0)[0] : AIC=-1788.084, Time=0.12 sec


Best model: ARIMA(3,1,0)(0,0,0)[0] intercept
Total fit time: 2.411 seconds
===== SARIMAX Results =====
Dep. Variable: y No. Observations: 449
Model: SARIMAX(3, 1, 0) Log Likelihood: 900.135
Date: Fri, 28 Nov 2025 AIC: -1790.270
Time: 02:56:26 BIC: -1769.746
Sample: 0 HQIC: -1782.179
Covariance Type: opg
===== coef std err z P>|z| [0.025 0.975]
intercept 0.0032 0.002 2.041 0.041 0.000 0.006
ar.L1 -0.0888 0.039 -2.267 0.023 -0.166 -0.012
ar.L2 0.0793 0.045 1.781 0.075 -0.008 0.166
ar.L3 -0.0720 0.046 -1.550 0.121 -0.163 0.019
sigma2 0.0011 4.16e-05 25.306 0.000 0.001 0.001
=====
Ljung-Box (L1) (Q): 0.01 Jarque-Bera (JB): 405.46
Prob(Q): 0.93 Prob(JB): 0.00
Heteroskedasticity (H): 0.93 Skew: -0.20
Prob(H) (two-sided): 0.65 Kurtosis: 7.64
=====
```



```
model_autoARIMA = auto_arima(train_data, test='adf', max_p=3, max_q=3, seasonal=False)
```

PREDICTING STOCK PRICES

ARIMA :



forecasted about 60 days

Configuration

Model Loading

Choose method:

- Upload File
- Enter Path

Upload your .pkl model

Drag and drop file here
Limit 200MB per file • PKL

[Browse files](#)

arima_stock_model.pkl 1.6MB [X](#)

Forecast Settings

Forecast Days

Fetch Recent Stock Data [?](#)

[Generate Predictions](#)

Built with Streamlit & ARIMA

Powered by pmdarima & yfinance

ARIMA Stock Price Prediction App

Predict future stock prices using your trained ARIMA model

Model loaded successfully!

Stock Symbol	ARIMA Order	Training Date	Forecast Period
NVDA	(1,1,0)	2025-11-24	16 days

Fetched 126 days of recent data

Historical Data Predictions Export & Info

Price Predictions

Model trained on log-transformed data - predictions converted to actual prices

Current Price	Day 1 Prediction	Day 16 Prediction
\$182.55	\$135.59 ↓ -25.73%	\$139.83 ↓ -23.40%

NVDA Price Forecast - Next 16 Days

The chart displays the historical price of NVDA from approximately November 2023 to early December 2023. A vertical dashed line at the end of the historical data marks the 'Forecast Start'. The 'Historical' data is shown as a blue line, the 'Forecast' as a red line, and the 'Forecast Start' as a vertical dotted line.

THE OPPORTUNITY

“No model can perfectly predict stock prices because markets are influenced by external events like news, earnings reports, and macroeconomic factors.

However, for a statistical baseline model, this ARIMA approach performs well and is appropriate for short-term forecasting.”

Key learning 1

Data Preparation is Critical because Log transformation helped stabilize variance and Raw stock prices contain trend and volatility. BE CAREFUL WHILE CLEANING

Key learning 2

ARIMA is a Strong Baseline for Short-Term Forecasting

Key learning 3

Model Logic Must Match Business Context

Thank you